

Airbnb

February 24, 2025

Part A: Airbnb Price Prediction and Insights

Deliverables

1. Data Exploration and Preprocessing

```
[241]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.simplefilter(action='ignore')

# Load the dataset
df = pd.read_csv("Airbnb_data.csv")

# Display the first 5 rows
df.head()
```

```
[241]:
```

	id	log_price	property_type	room_type	\
0	6901257	5.010635	Apartment	Entire home/apt	
1	6304928	5.129899	Apartment	Entire home/apt	
2	7919400	4.976734	Apartment	Entire home/apt	
3	13418779	6.620073	House	Entire home/apt	
4	3808709	4.744932	Apartment	Entire home/apt	

	amenities	accommodates	bathrooms	\
0	{"Wireless Internet","Air conditioning",Kitche...	3	1.0	
1	{"Wireless Internet","Air conditioning",Kitche...	7	1.0	
2	{TV,"Cable TV","Wireless Internet","Air condit...	5	1.0	
3	{TV,"Cable TV",Internet,"Wireless Internet",Ki...	4	1.0	
4	{TV,Internet,"Wireless Internet","Air conditio...	2	1.0	

	bed_type	cancellation_policy	cleaning_fee	...	latitude	longitude	\
0	Real Bed	strict	True	...	40.696524	-73.991617	
1	Real Bed	strict	True	...	40.766115	-73.989040	
2	Real Bed	moderate	True	...	40.808110	-73.943756	
3	Real Bed	flexible	True	...	37.772004	-122.431619	
4	Real Bed	moderate	True	...	38.925627	-77.034596	

	name	neighbourhood	
0	Beautiful brownstone 1-bedroom	Brooklyn Heights	
1	Superb 3BR Apt Located Near Times Square	Hell's Kitchen	
2	The Garden Oasis	Harlem	
3	Beautiful Flat in the Heart of SF!	Lower Haight	
4	Great studio in midtown DC	Columbia Heights	

	number_of_reviews	review_scores_rating	
0	2	100.0	
1	6	93.0	
2	10	92.0	
3	0	NaN	
4	4	40.0	

	thumbnail_url	zipcode	bedrooms	beds
0	https://a0.muscache.com/im/pictures/6d7cbbf7-c...	11201	1.0	1.0
1	https://a0.muscache.com/im/pictures/348a55fe-4...	10019	3.0	3.0
2	https://a0.muscache.com/im/pictures/6fae5362-9...	10027	1.0	3.0
3	https://a0.muscache.com/im/pictures/72208dad-9...	94117	2.0	2.0
4		NaN	20009	0.0

[5 rows x 29 columns]

```
[243]: # Check dataset info
df.info()

# Check summary statistics for numerical columns
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74111 entries, 0 to 74110
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     74111 non-null  int64
1   log_price              74111 non-null  float64
2   property_type          74111 non-null  object
3   room_type              74111 non-null  object
4   amenities               74111 non-null  object
5   accommodates            74111 non-null  int64
6   bathrooms              73911 non-null  float64
7   bed_type               74111 non-null  object
8   cancellation_policy     74111 non-null  object
9   cleaning_fee           74111 non-null  bool
10  city                   74111 non-null  object
11  description             74111 non-null  object
12  first_review           58247 non-null  object
```

```

13 host_has_profile_pic      73923 non-null object
14 host_identity_verified    73923 non-null object
15 host_response_rate        55812 non-null object
16 host_since                73923 non-null object
17 instant_bookable          74111 non-null object
18 last_review               58284 non-null object
19 latitude                  74111 non-null float64
20 longitude                 74111 non-null float64
21 name                      74111 non-null object
22 neighbourhood             67239 non-null object
23 number_of_reviews         74111 non-null int64
24 review_scores_rating      57389 non-null float64
25 thumbnail_url            65895 non-null object
26 zipcode                   73143 non-null object
27 bedrooms                  74020 non-null float64
28 beds                     73980 non-null float64

```

dtypes: bool(1), float64(7), int64(3), object(18)

memory usage: 15.9+ MB

```

[243]:
      id      log_price  accommodates  bathrooms  latitude \
count  7.411100e+04  74111.000000  74111.000000  73911.000000  74111.000000
mean    1.126662e+07    4.782069    3.155146    1.235263    38.445958
std     6.081735e+06    0.717394    2.153589    0.582044    3.080167
min     3.440000e+02    0.000000    1.000000    0.000000    33.338905
25%     6.261964e+06    4.317488    2.000000    1.000000    34.127908
50%     1.225415e+07    4.709530    2.000000    1.000000    40.662138
75%     1.640226e+07    5.220356    4.000000    1.000000    40.746096
max     2.123090e+07    7.600402   16.000000    8.000000    42.390437

      longitude  number_of_reviews  review_scores_rating  bedrooms \
count  74111.000000    74111.000000    57389.000000  74020.000000
mean    -92.397525    20.900568    94.067365    1.265793
std     21.705322    37.828641    7.836556    0.852143
min    -122.511500    0.000000    20.000000    0.000000
25%    -118.342374    1.000000    92.000000    1.000000
50%    -76.996965    6.000000    96.000000    1.000000
75%    -73.954660    23.000000   100.000000    1.000000
max    -70.985047   605.000000   100.000000   10.000000

      beds
count  73980.000000
mean    1.710868
std     1.254142
min     0.000000
25%     1.000000
50%     1.000000
75%     2.000000

```

```
max      18.000000
```

```
[245]: # Check number of missing values in each column
df.isnull().sum()
```

```
[245]: id      0
log_price    0
property_type 0
room_type    0
amenities    0
accommodates 0
bathrooms    200
bed_type     0
cancellation_policy 0
cleaning_fee 0
city         0
description   0
first_review  15864
host_has_profile_pic 188
host_identity_verified 188
host_response_rate 18299
host_since    188
instant_bookable 0
last_review   15827
latitude      0
longitude     0
name          0
neighbourhood 6872
number_of_reviews 0
review_scores_rating 16722
thumbnail_url 8216
zipcode       968
bedrooms      91
beds          131
dtype: int64
```

```
[247]: # Fill missing values for numerical columns with median
num_cols = ['bathrooms', 'bedrooms', 'beds', 'review_scores_rating']
df[num_cols] = df[num_cols].fillna(df[num_cols].median())
# Verify changes
df[num_cols].isnull().sum()
```

```
[247]: bathrooms    0
bedrooms       0
beds           0
review_scores_rating 0
dtype: int64
```

```
[249]: # Convert non-null values to string first, then remove '%' and convert to float
df['host_response_rate'] = df['host_response_rate'].astype(str).str.rstrip('%')

# Convert column to numeric, setting errors='coerce' to handle any unexpected
↳ values
df['host_response_rate'] = pd.to_numeric(df['host_response_rate'],
↳ errors='coerce')

# Fill missing values with median
df['host_response_rate'].fillna(df['host_response_rate'].median(), inplace=True)

# Verify changes
df['host_response_rate'].isnull().sum()
```

[249]: 0

```
[251]: # Convert dates to datetime format
df['host_since'] = pd.to_datetime(df['host_since'])
df['first_review'] = pd.to_datetime(df['first_review'])
df['last_review'] = pd.to_datetime(df['last_review'])

# Fill missing dates with the most common date
df['host_since'].fillna(df['host_since'].mode()[0], inplace=True)
df['first_review'].fillna(df['first_review'].mode()[0], inplace=True)
df['last_review'].fillna(df['last_review'].mode()[0], inplace=True)

# Create a new feature: How long the host has been on Airbnb
df['host_years'] = 2025 - df['host_since'].dt.year

# Verify changes
df[['host_since', 'first_review', 'last_review', 'host_years']].isnull().sum()
```

```
[251]: host_since      0
first_review      0
last_review       0
host_years        0
dtype: int64
```

```
[253]: # List of boolean columns
bool_cols = ['host_has_profile_pic', 'host_identity_verified',
↳ 'instant_bookable']

# Convert 't' to True and 'f' to False
for col in bool_cols:
    df[col] = df[col].map({'t': True, 'f': False})

# Fill missing values with False and convert to integer (0 or 1)
```

```
df[bool_cols] = df[bool_cols].fillna(False).astype(int)

# Verify changes
df[bool_cols].isnull().sum()
```

```
[253]: host_has_profile_pic      0
      host_identity_verified    0
      instant_bookable         0
      dtype: int64
```

```
[255]: df['neighbourhood'].fillna("Unknown", inplace=True) # Fill with "Unknown"
      df['zipcode'].fillna(df['zipcode'].mode()[0], inplace=True) #Fill with mode

# Verify changes
df[['neighbourhood', 'zipcode']].isnull().sum()
```

```
[255]: neighbourhood      0
      zipcode            0
      dtype: int64
```

```
[257]: # Dropping some columns which are not much useful
      df.drop(columns=['thumbnail_url', 'description'], inplace=True)

# Verify remaining missing values
df.isnull().sum()
```

```
[257]: id                      0
      log_price                0
      property_type            0
      room_type                0
      amenities                0
      accommodates             0
      bathrooms                0
      bed_type                 0
      cancellation_policy      0
      cleaning_fee             0
      city                     0
      first_review             0
      host_has_profile_pic     0
      host_identity_verified    0
      host_response_rate       0
      host_since               0
      instant_bookable         0
      last_review              0
      latitude                 0
      longitude                0
      name                     0
```

```

neighbourhood          0
number_of_reviews      0
review_scores_rating    0
zipcode                0
bedrooms               0
beds                   0
host_years              0
dtype: int64

```

```

[259]: # Define function to remove outliers using IQR
def remove_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

# Apply to numerical columns with potential outliers
df = remove_outliers(df, "log_price")
df = remove_outliers(df, "accommodates")
df = remove_outliers(df, "bathrooms")
df = remove_outliers(df, "bedrooms")
df = remove_outliers(df, "beds")

# Verify shape after removing outliers
df.shape

```

[259]: (34576, 28)

```

[261]: #Feature Engineering such as neighborhood popularity, number of amenities, and
        ↪ host activity metrics
# Compute the number of listings per neighborhood
neighborhood_counts = df['neighbourhood'].value_counts()

# Compute the average review score per neighborhood
neighborhood_avg_review = df.groupby('neighbourhood')['review_scores_rating'].
        ↪ mean()

# Normalize and combine these factors into a popularity score
df['neighborhood_popularity'] = df['neighbourhood'].map(lambda x:
        ↪ (neighborhood_counts[x] - neighborhood_counts.min())
                                                    / (neighborhood_counts.
        ↪ max() - neighborhood_counts.min()))
df['neighborhood_avg_review'] = df['neighbourhood'].map(lambda x:
        ↪ (neighborhood_avg_review[x] - neighborhood_avg_review.min())

```

```

/
↪(neighborhood_avg_review.max() - neighborhood_avg_review.min()))

# Weighted score: 70% listings count + 30% average review score
df['neighborhood_popularity'] = (df['neighborhood_popularity'] * 0.7) +
↪(df['neighborhood_avg_review'] * 0.3)

# Fill NaNs with the mean value
df['neighborhood_popularity'].fillna(df['neighborhood_popularity'].mean(),
↪inplace=True)

# Verify the new feature
df[['neighbourhood', 'neighborhood_popularity']].head()

```

```

[261]:      neighbourhood  neighborhood_popularity
0    Brooklyn Heights          0.247012
5         Noe Valley          0.256796
6           Unknown          0.906962
7         Downtown          0.248859
8  Richmond District          0.243277

```

```

[262]: # Create a new column for number of amenities
df['num_amenities'] = df['amenities'].apply(lambda x: len(x.split(',')) if
↪isinstance(x, str) else 0)

# Verify the new feature
df[['amenities', 'num_amenities']].head()

```

```

[262]:      amenities  num_amenities
0  {"Wireless Internet","Air conditioning",Kitche...      9
5  {"TV","Wireless Internet",Heating,"Smoke detecto...     10
6  {"TV,Internet","Wireless Internet","Air conditio...     21
7  {"TV","Cable TV","Wireless Internet","Wheelchair...     26
8  {"TV","Cable TV","Wireless Internet","Pets live ...     21

```

```

[265]: # Convert 'host_response_rate' from percentage string to numeric
df['host_response_rate'] = df['host_response_rate'].astype(str).str.rstrip('%').
↪astype(float) / 100

# Create a feature for the number of days the host has been on Airbnb
df['host_since'] = pd.to_datetime(df['host_since'])
df['host_days_active'] = (pd.to_datetime('today') - df['host_since']).dt.days

# Verify the new features
df[['host_response_rate', 'host_days_active']].head()

```



```
[265]:    host_response_rate  host_days_active
0                1.0           4718
5                1.0           2819
6                1.0           2915
7                1.0           4300
8                1.0           3553
```

2. Model Development

```
[268]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Convert 'amenities' to numeric (count of amenities)
df['num_amenities'] = df['amenities'].apply(lambda x: len(x.split(',')) if
↳ isinstance(x, str) else 0)
df = df.drop(columns=['amenities'])

# Convert categorical features using one-hot encoding
df = pd.get_dummies(df, columns=['city', 'neighbourhood', 'zipcode'],
↳ drop_first=True)

# Convert boolean-like columns ('t' and 'f') to numeric (1 and 0)
bool_cols = ['host_has_profile_pic', 'host_identity_verified',
↳ 'instant_bookable']
df[bool_cols] = df[bool_cols].replace({'t': 1, 'f': 0})

# Define target variable (log-transformed price)
y = df['log_price']
X = df.drop(columns=['log_price', 'id', 'name']) # Drop non-relevant columns

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[270]: X_train = pd.get_dummies(X_train, drop_first=True)
X_test = pd.get_dummies(X_test, drop_first=True)
```

```
[90]: # Ensure X_train and X_test have the same columns
X_train, X_test = X_train.align(X_test, join='left', axis=1, fill_value=0)
```

```
[38]: # Linear Regression Model
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
y_pred_lr = lr_model.predict(X_test)
```

```
# Random Forest Model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
y_pred_rf = rf_model.predict(X_test)
```

3. Model Evaluation

```
[42]: def evaluate_model(y_test, y_pred, model_name):
    print(f" {model_name} Performance:")
    print(f" RMSE: {np.sqrt(mean_squared_error(y_test, y_pred)):.2f}")
    print(f" MAE: {mean_absolute_error(y_test, y_pred):.2f}")
    print(f" R2 Score: {r2_score(y_test, y_pred):.2f}")
    print("-" * 40)

# Evaluate both models
evaluate_model(y_test, y_pred_lr, "Linear Regression")
evaluate_model(y_test, y_pred_rf, "Random Forest Regressor")
```

Linear Regression Performance:

RMSE: 0.54

MAE: 0.43

R² Score: 0.02

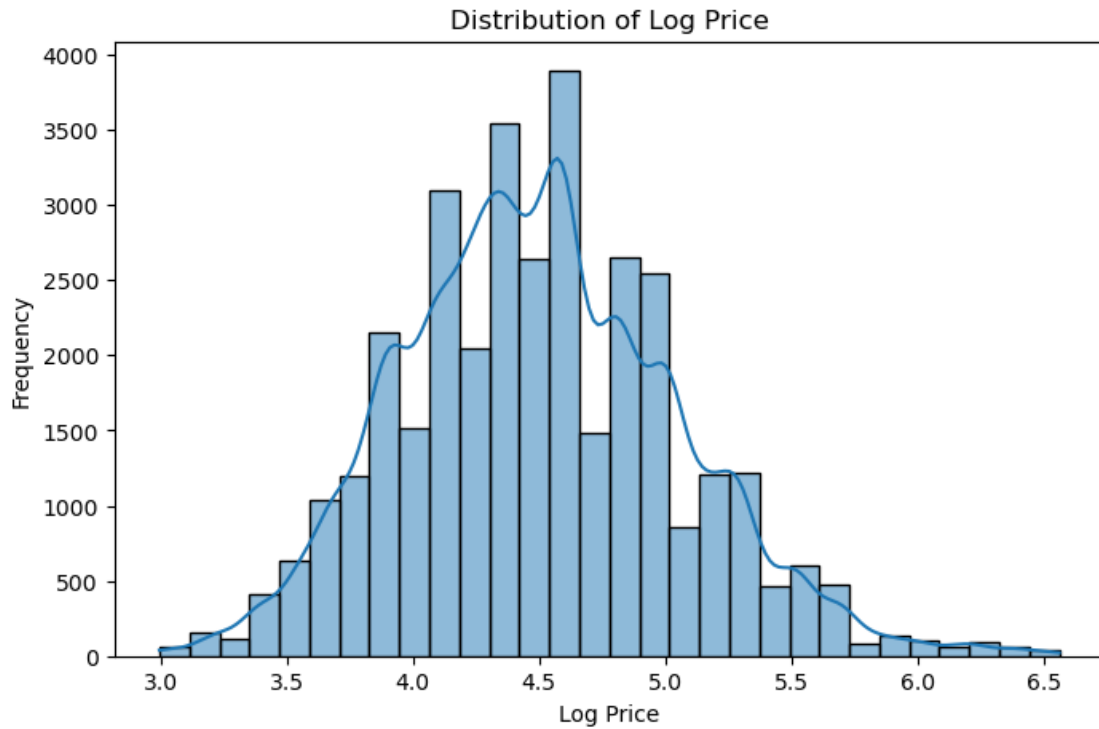
Random Forest Regressor Performance:

RMSE: 0.37

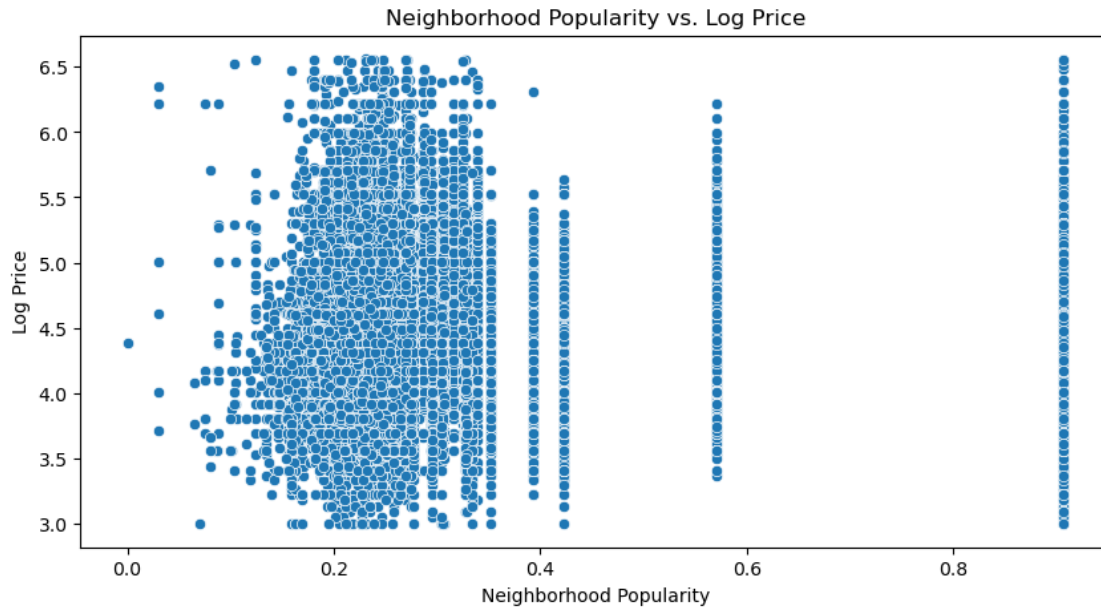
MAE: 0.28

R² Score: 0.54

```
[66]: # Visualizations
plt.figure(figsize=(8, 5))
sns.histplot(df['log_price'], bins=30, kde=True)
plt.title("Distribution of Log Price")
plt.xlabel("Log Price")
plt.ylabel("Frequency")
plt.show()
```



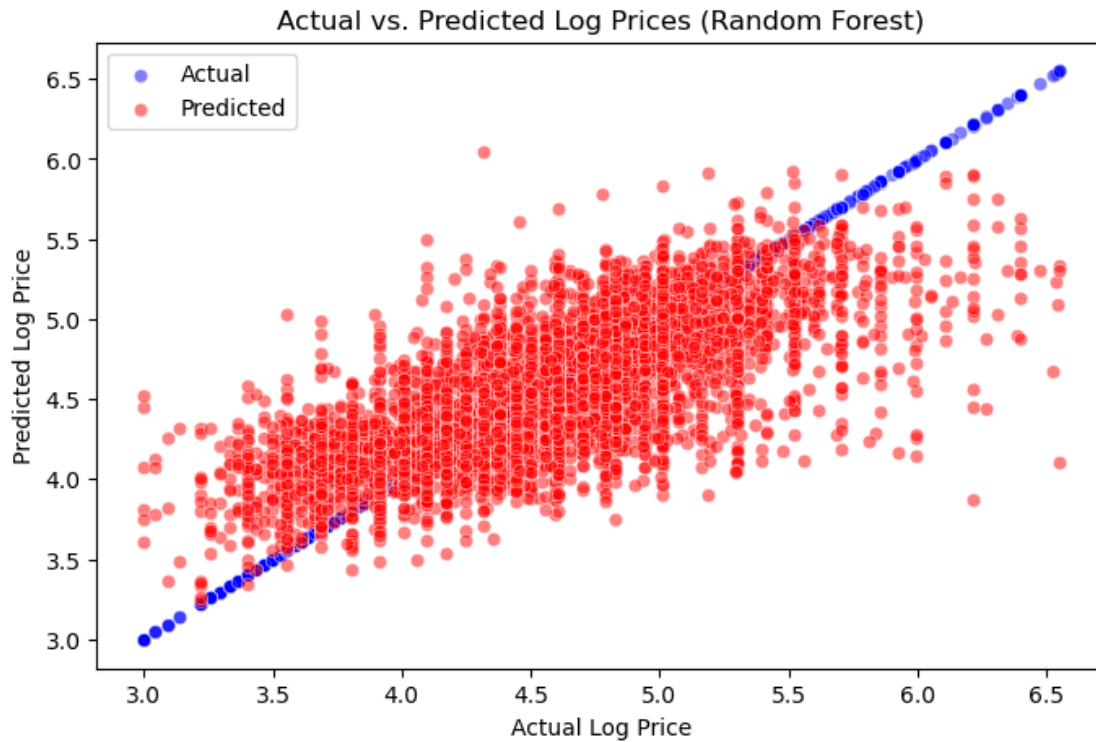
```
[68]: plt.figure(figsize=(10, 5))
sns.scatterplot(x=df['neighborhood_popularity'], y=df['log_price'])
plt.title("Neighborhood Popularity vs. Log Price")
plt.xlabel("Neighborhood Popularity")
plt.ylabel("Log Price")
plt.show()
```



```
[272]: plt.figure(figsize=(8, 5))

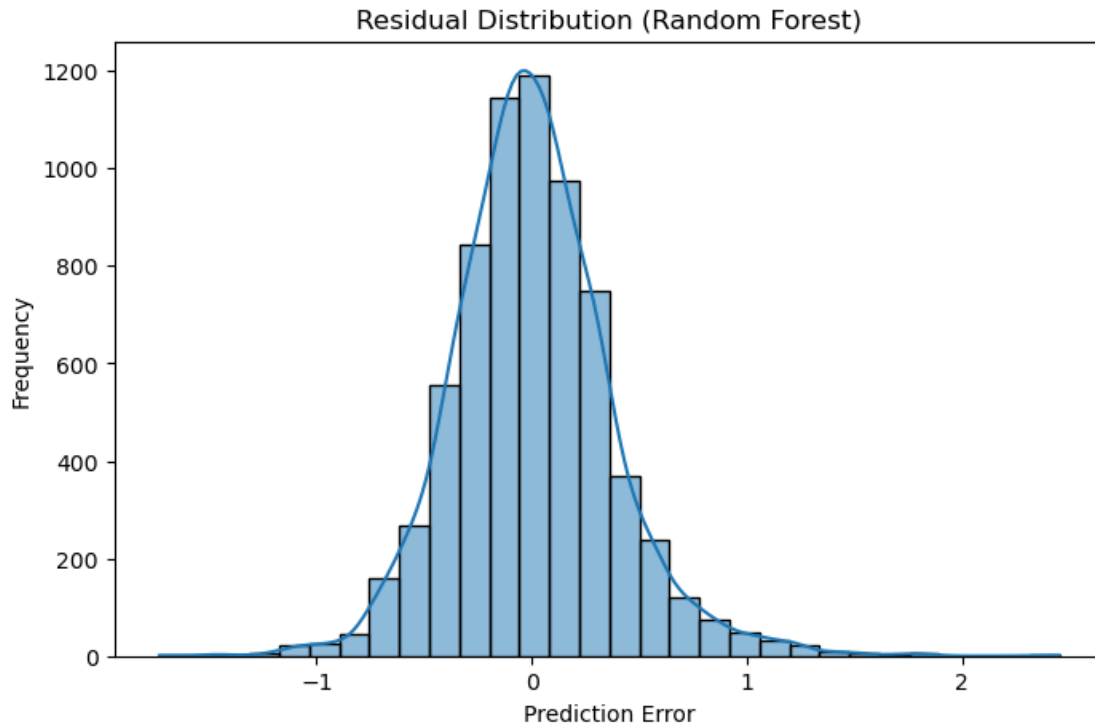
# Plot actual values in one color and predicted values in another
sns.scatterplot(x=y_test, y=y_test, color="blue", alpha=0.5, label="Actual") # Blue for actual values
sns.scatterplot(x=y_test, y=y_pred_rf, color="red", alpha=0.5, label="Predicted") # Red for predicted values

plt.title("Actual vs. Predicted Log Prices (Random Forest)")
plt.xlabel("Actual Log Price")
plt.ylabel("Predicted Log Price")
plt.legend() # Show the legend
plt.show()
```



```
[74]: residuals = y_test - y_pred_rf

plt.figure(figsize=(8, 5))
sns.histplot(residuals, bins=30, kde=True)
plt.title("Residual Distribution (Random Forest)")
plt.xlabel("Prediction Error")
plt.ylabel("Frequency")
plt.show()
```



```
[86]: # Tuning the Model
from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [None, 10, 30],
    'min_samples_split': [2, 5, 10]
}
grid_search = GridSearchCV(RandomForestRegressor(n_estimators=100,
    random_state=42, n_jobs=-1),
                           param_grid, cv=3, scoring='r2', n_jobs=-1)
grid_search.fit(X_train, y_train)

best_rf = grid_search.best_estimator_
y_pred_best_rf = best_rf.predict(X_test)

evaluate_model(y_test, y_pred_best_rf, "Grid Search Tuned RF")
```

Grid Search Tuned RF Performance:

RMSE: 0.37

MAE: 0.28

R² Score: 0.54

```
[98]: from xgboost import XGBRegressor

# Define XGBoost model with default parameters
xgb = XGBRegressor(random_state=42)

# Train the model
xgb.fit(X_train, y_train)

# Predict on test data
y_pred_xgb = xgb.predict(X_test)

# Evaluate XGBoost
evaluate_model(y_test, y_pred_xgb, "XGBoost Regressor")
```

XGBoost Regressor Performance:
 RMSE: 0.36
 MAE: 0.27
 R² Score: 0.56

```
[276]: # Insights about the factors influencing Airbnb prices

# Create a DataFrame of feature importance
feature_importance = pd.DataFrame({'Feature': X_train.columns, 'Importance':
    ↳xgb.feature_importances_})

# Sort by importance and show top 10
top_10_features = feature_importance.sort_values(by='Importance',
    ↳ascending=False).head(10)
plt.figure(figsize=(10, 5))
plt.barh(top_10_features['Feature'], top_10_features['Importance'],
    ↳color='skyblue')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.title('Top 10 most important factors influencing Airbnb prices')
plt.gca().invert_yaxis() # Invert y-axis to show the highest at the top
plt.show()
```

