

Customer Churn Prediction

Part B: Customer Churn Prediction

Deliverables

1. A data exploration and preprocessing notebook or report that analyzes the dataset, handles missing values, and prepares the data for modeling.

```
[82]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
import warnings
warnings.simplefilter(action='ignore')

# Load the dataset
file_path = "customer_data.csv"
df = pd.read_csv(file_path)

# Display basic information about the dataset
df.info()
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7043 entries, 0 to 7042
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64
6	PhoneService	7043 non-null	object
7	MultipleLines	7043 non-null	object
8	InternetService	7043 non-null	object
9	OnlineSecurity	7043 non-null	object
10	OnlineBackup	7043 non-null	object
11	DeviceProtection	7043 non-null	object

```

12 TechSupport      7043 non-null  object
13 StreamingTV      7043 non-null  object
14 StreamingMovies  7043 non-null  object
15 Contract         7043 non-null  object
16 PaperlessBilling 7043 non-null  object
17 PaymentMethod    7043 non-null  object
18 MonthlyCharges   7043 non-null  float64
19 TotalCharges     7032 non-null  float64
20 Churn            7043 non-null  object

```

dtypes: float64(2), int64(2), object(17)

memory usage: 1.1+ MB

```

[82]:  customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0  7590-VHVEG  Female                0      Yes           No         1           No
1  5575-GNVDE   Male                0      No           No        34           Yes
2  3668-QPYBK   Male                0      No           No         2           Yes
3  7795-CFOCW   Male                0      No           No        45           No
4  9237-HQITU   Female                0      No           No         2           Yes

```

```

      MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection  \
0  No phone service                DSL                No  ...                No
1                    No                DSL                Yes  ...                Yes
2                    No                DSL                Yes  ...                No
3  No phone service                DSL                Yes  ...                Yes
4                    No      Fiber optic                No  ...                No

```

```

      TechSupport  StreamingTV  StreamingMovies      Contract  PaperlessBilling  \
0              No              No              No  Month-to-month              Yes
1              No              No              No      One year              No
2              No              No              No  Month-to-month              Yes
3              Yes              No              No      One year              No
4              No              No              No  Month-to-month              Yes

```

```

      PaymentMethod  MonthlyCharges  TotalCharges  Churn
0      Electronic check           29.85          29.85   No
1          Mailed check           56.95         1889.50   No
2          Mailed check           53.85          108.15  Yes
3  Bank transfer (automatic)         42.30         1840.75   No
4      Electronic check           70.70          151.65  Yes

```

[5 rows x 21 columns]

```
[84]: df.describe()
```

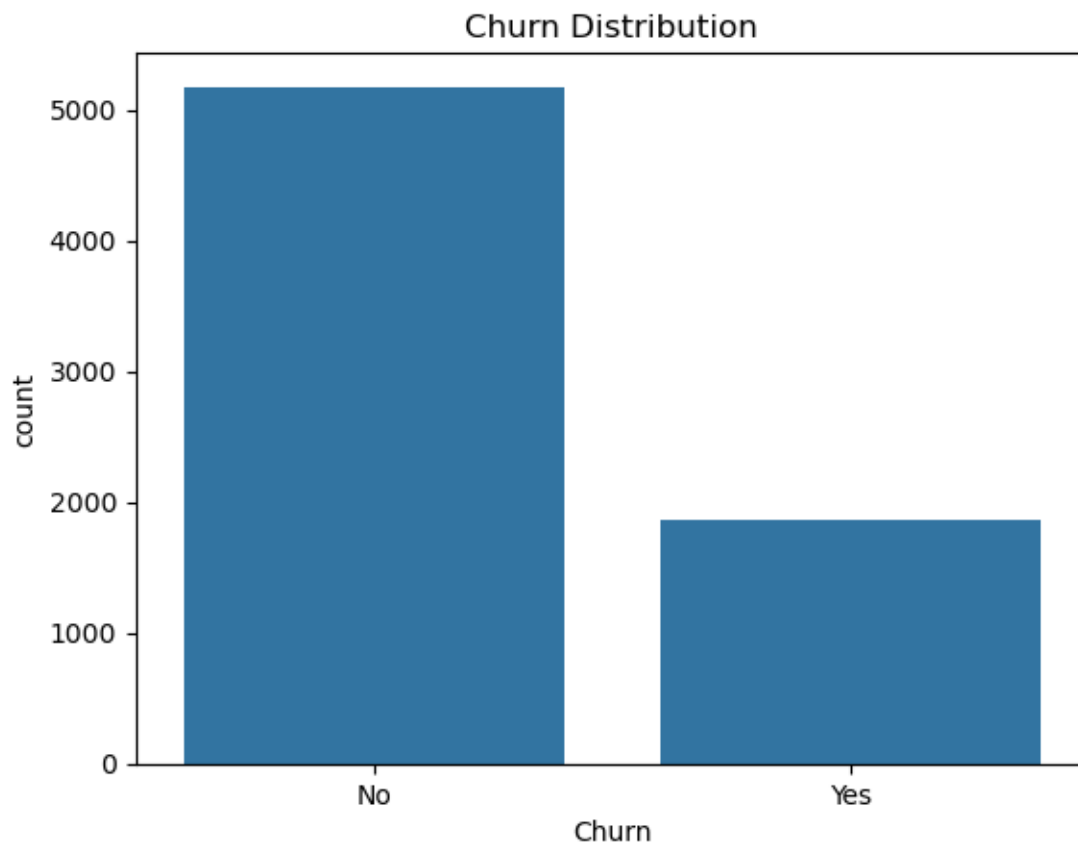
```

[84]:      SeniorCitizen      tenure  MonthlyCharges  TotalCharges
count      7043.000000  7043.000000    7043.000000    7032.000000
mean         0.162147    32.371149      64.761692    2283.300441
std         0.368612    24.559481     30.090047    2266.771362

```

min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

```
[86]: # Visualizing Churn distribution
sns.countplot(x='Churn', data=df)
plt.title('Churn Distribution')
plt.show()
```



```
[88]: df.isnull().sum()
```

```
[88]: customerID      0
gender              0
SeniorCitizen      0
Partner            0
Dependents         0
tenure             0
PhoneService       0
```

```

MultipleLines      0
InternetService    0
OnlineSecurity     0
OnlineBackup       0
DeviceProtection   0
TechSupport        0
StreamingTV        0
StreamingMovies    0
Contract           0
PaperlessBilling   0
PaymentMethod      0
MonthlyCharges     0
TotalCharges       11
Churn              0
dtype: int64

```

```

[90]: # Handle missing values in TotalCharges by filling them with the median value
df["TotalCharges"].fillna(df["TotalCharges"].median(), inplace=True)

# Drop the customerID column as it is not useful for prediction
df.drop(columns=["customerID"], inplace=True)

# Convert categorical variables with "Yes/No" values to binary (1/0)
yes_no_cols = ["Partner", "Dependents", "PhoneService", "PaperlessBilling",
               ↪ "Churn"]
for col in yes_no_cols:
    df[col] = df[col].map({"Yes": 1, "No": 0})

# Convert gender to binary (Male=1, Female=0)
df["gender"] = df["gender"].map({"Male": 1, "Female": 0})

# Keep SeniorCitizen as a numerical feature
df["SeniorCitizen"] = df["SeniorCitizen"].astype(int)

# Identify categorical variables with more than two unique values
multi_category_cols = ["Contract", "InternetService", "PaymentMethod",
                       ↪ "MultipleLines", "OnlineSecurity",
                           "OnlineBackup", "DeviceProtection", "TechSupport",
                       ↪ "StreamingTV", "StreamingMovies"]

# Apply one-hot encoding to multi-category variables
df = pd.get_dummies(df, columns=multi_category_cols, drop_first=True)

# Scale numerical features
scaler = StandardScaler()
df[["tenure", "MonthlyCharges", "TotalCharges"]] = scaler.
    ↪ fit_transform(df[["tenure", "MonthlyCharges", "TotalCharges"]])

```

```
# Display the processed dataset info
df.info()
df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7043 entries, 0 to 7042
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	gender	7043 non-null	int64
1	SeniorCitizen	7043 non-null	int32
2	Partner	7043 non-null	int64
3	Dependents	7043 non-null	int64
4	tenure	7043 non-null	float64
5	PhoneService	7043 non-null	int64
6	PaperlessBilling	7043 non-null	int64
7	MonthlyCharges	7043 non-null	float64
8	TotalCharges	7043 non-null	float64
9	Churn	7043 non-null	int64
10	Contract_One year	7043 non-null	bool
11	Contract_Two year	7043 non-null	bool
12	InternetService_Fiber optic	7043 non-null	bool
13	InternetService_No	7043 non-null	bool
14	PaymentMethod_Credit card (automatic)	7043 non-null	bool
15	PaymentMethod_Electronic check	7043 non-null	bool
16	PaymentMethod_Mailed check	7043 non-null	bool
17	MultipleLines_No phone service	7043 non-null	bool
18	MultipleLines_Yes	7043 non-null	bool
19	OnlineSecurity_No internet service	7043 non-null	bool
20	OnlineSecurity_Yes	7043 non-null	bool
21	OnlineBackup_No internet service	7043 non-null	bool
22	OnlineBackup_Yes	7043 non-null	bool
23	DeviceProtection_No internet service	7043 non-null	bool
24	DeviceProtection_Yes	7043 non-null	bool
25	TechSupport_No internet service	7043 non-null	bool
26	TechSupport_Yes	7043 non-null	bool
27	StreamingTV_No internet service	7043 non-null	bool
28	StreamingTV_Yes	7043 non-null	bool
29	StreamingMovies_No internet service	7043 non-null	bool
30	StreamingMovies_Yes	7043 non-null	bool

```
dtypes: bool(21), float64(3), int32(1), int64(6)
```

```
memory usage: 667.3 KB
```

```
[90]:   gender  SeniorCitizen  Partner  Dependents   tenure  PhoneService  \
0         0              0         1         0 -1.277445             0
1         1              0         0         0  0.066327             1
```

2	1	0	0	0 -1.236724	1
3	1	0	0	0 0.514251	0
4	0	0	0	0 -1.236724	1

	PaperlessBilling	MonthlyCharges	TotalCharges	Churn	...	\
0	1	-1.160323	-0.994242	0	...	
1	0	-0.259629	-0.173244	0	...	
2	1	-0.362660	-0.959674	1	...	
3	0	-0.746535	-0.194766	0	...	
4	1	0.197365	-0.940470	1	...	

	OnlineBackup_No internet service	OnlineBackup_Yes	\
0	False	True	
1	False	False	
2	False	True	
3	False	False	
4	False	False	

	DeviceProtection_No internet service	DeviceProtection_Yes	\
0	False	False	
1	False	True	
2	False	False	
3	False	True	
4	False	False	

	TechSupport_No internet service	TechSupport_Yes	\
0	False	False	
1	False	False	
2	False	False	
3	False	True	
4	False	False	

	StreamingTV_No internet service	StreamingTV_Yes	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	StreamingMovies_No internet service	StreamingMovies_Yes
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

[5 rows x 31 columns]

2. A machine learning model capable of predicting customer churn.

```
[93]: from sklearn.model_selection import train_test_split, GridSearchCV
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score, classification_report, \
          precision_score, recall_score, f1_score

      # Split the data into features and target variable
      X = df.drop(columns=["Churn"])
      y = df["Churn"]

      # Split into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, \
          random_state=42)

      # Train a Random Forest model
      model = RandomForestClassifier(n_estimators=100, random_state=42)
      model.fit(X_train, y_train)

      # Make predictions
      y_pred = model.predict(X_test)
```

3. An evaluation of model performance using appropriate metrics (such as accuracy, precision, recall, F1 score, etc.)

```
[96]: # Evaluate the model
      accuracy = accuracy_score(y_test, y_pred)
      precision = precision_score(y_test, y_pred)
      recall = recall_score(y_test, y_pred)
      f1 = f1_score(y_test, y_pred)

      print("Accuracy:", accuracy)
      print("Precision:", precision)
      print("Recall:", recall)
      print("F1 Score:", f1)
      print("Classification Report:")
      print(classification_report(y_test, y_pred))
```

Accuracy: 0.7934705464868701

Precision: 0.6553030303030303

Recall: 0.46380697050938335

F1 Score: 0.543171114599686

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.66	0.46	0.54	373

accuracy			0.79	1409
macro avg	0.74	0.69	0.70	1409
weighted avg	0.78	0.79	0.78	1409

```
[98]: # Hyperparameter tuning using GridSearchCV
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(RandomForestClassifier(random_state=42), param_grid,
    cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

# Best model
best_model = grid_search.best_estimator_

# Make predictions
y_pred_best = best_model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred_best)
precision = precision_score(y_test, y_pred_best)
recall = recall_score(y_test, y_pred_best)
f1 = f1_score(y_test, y_pred_best)

print("Best Model Parameters:", grid_search.best_params_)
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Classification Report:")
print(classification_report(y_test, y_pred))
```


```
Best Model Parameters: {'max_depth': 10, 'min_samples_split': 10,
'n_estimators': 150}
Accuracy: 0.808374733853797
Precision: 0.6872727272727273
Recall: 0.5067024128686327
F1 Score: 0.5833333333333334
Classification Report:
```

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1036
1	0.66	0.46	0.54	373

accuracy			0.79	1409
macro avg	0.74	0.69	0.70	1409
weighted avg	0.78	0.79	0.78	1409

```
[100]: # Feature importance analysis
feature_importances = pd.DataFrame({'Feature': X.columns, 'Importance':
    ↳best_model.feature_importances_})
feature_importances = feature_importances.sort_values(by='Importance',
    ↳ascending=False)

plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importances)
plt.title('Feature Importance')
plt.show()
```



```
[114]: # Function to predict churn for new data
def predict_churn(new_data):
    new_data = pd.DataFrame(new_data)
    new_data[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.
    ↳transform(new_data[['tenure', 'MonthlyCharges', 'TotalCharges']])
    return best_model.predict(new_data)

# Example new customer data
new_customer = {
    'gender': [1],
    'SeniorCitizen': [0],
    'Partner': [1],
    'Dependents': [0],
```

```

'tenure': [12],
'PhoneService': [1],
'PaperlessBilling': [1],
'MonthlyCharges': [50.0],
'TotalCharges': [600.0],
'Contract_One year': [0],
'Contract_Two year': [1],
'InternetService_Fiber optic': [1],
'InternetService_No': [0],
'PaymentMethod_Credit card (automatic)': [0],
'PaymentMethod_Electronic check': [1],
'PaymentMethod_Mailed check': [0],
'MultipleLines_No phone service': [0],
'MultipleLines_Yes': [1],
'OnlineSecurity_No internet service': [0],
'OnlineSecurity_Yes': [1],
'OnlineBackup_No internet service': [0],
'OnlineBackup_Yes': [1],
'DeviceProtection_No internet service': [0],
'DeviceProtection_Yes': [1],
'TechSupport_No internet service': [0],
'TechSupport_Yes': [1],
'StreamingTV_No internet service': [0],
'StreamingTV_Yes': [1],
'StreamingMovies_No internet service': [0],
'StreamingMovies_Yes': [1]
}

# Convert to DataFrame and predict
new_customer_df = pd.DataFrame(new_customer)
prediction = predict_churn(new_customer_df)

# Print prediction result
print("Predicted Churn:", "Yes" if prediction[0] == 1 else "No")

```

Predicted Churn: No

— Interpretation and Actionable Insights — 1. High Monthly Charges and Low Tenure are strong indicators of churn. - Action: Offer discounts or loyalty programs to retain customers early. 2. Customers with No Internet Service are less likely to churn. - Action: Focus retention efforts on high-risk segments like fiber-optic users. 3. Paperless Billing and Electronic Check payment methods correlate with higher churn. - Action: Encourage automatic payments or credit card-based payments to reduce churn.

```

[117]: # Display the processed dataset info
df.info()
df.head()

```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 7043 entries, 0 to 7042

Data columns (total 31 columns):

#	Column	Non-Null Count	Dtype
0	gender	7043 non-null	int64
1	SeniorCitizen	7043 non-null	int32
2	Partner	7043 non-null	int64
3	Dependents	7043 non-null	int64
4	tenure	7043 non-null	float64
5	PhoneService	7043 non-null	int64
6	PaperlessBilling	7043 non-null	int64
7	MonthlyCharges	7043 non-null	float64
8	TotalCharges	7043 non-null	float64
9	Churn	7043 non-null	int64
10	Contract_One year	7043 non-null	bool
11	Contract_Two year	7043 non-null	bool
12	InternetService_Fiber optic	7043 non-null	bool
13	InternetService_No	7043 non-null	bool
14	PaymentMethod_Credit card (automatic)	7043 non-null	bool
15	PaymentMethod_Electronic check	7043 non-null	bool
16	PaymentMethod_Mailed check	7043 non-null	bool
17	MultipleLines_No phone service	7043 non-null	bool
18	MultipleLines_Yes	7043 non-null	bool
19	OnlineSecurity_No internet service	7043 non-null	bool
20	OnlineSecurity_Yes	7043 non-null	bool
21	OnlineBackup_No internet service	7043 non-null	bool
22	OnlineBackup_Yes	7043 non-null	bool
23	DeviceProtection_No internet service	7043 non-null	bool
24	DeviceProtection_Yes	7043 non-null	bool
25	TechSupport_No internet service	7043 non-null	bool
26	TechSupport_Yes	7043 non-null	bool
27	StreamingTV_No internet service	7043 non-null	bool
28	StreamingTV_Yes	7043 non-null	bool
29	StreamingMovies_No internet service	7043 non-null	bool
30	StreamingMovies_Yes	7043 non-null	bool

dtypes: bool(21), float64(3), int32(1), int64(6)

memory usage: 667.3 KB

```
[117]:  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService  \
0         0             0         1           0 -1.277445           0
1         1             0         0           0  0.066327           1
2         1             0         0           0 -1.236724           1
3         1             0         0           0  0.514251           0
4         0             0         0           0 -1.236724           1

      PaperlessBilling  MonthlyCharges  TotalCharges  Churn  ...  \
0                   1       -1.160323      -0.994242     0  ...
```

1	0	-0.259629	-0.173244	0	...
2	1	-0.362660	-0.959674	1	...
3	0	-0.746535	-0.194766	0	...
4	1	0.197365	-0.940470	1	...

	OnlineBackup_No	internet service	OnlineBackup_Yes	\
0		False	True	
1		False	False	
2		False	True	
3		False	False	
4		False	False	

	DeviceProtection_No	internet service	DeviceProtection_Yes	\
0		False	False	
1		False	True	
2		False	False	
3		False	True	
4		False	False	

	TechSupport_No	internet service	TechSupport_Yes	\
0		False	False	
1		False	False	
2		False	False	
3		False	True	
4		False	False	

	StreamingTV_No	internet service	StreamingTV_Yes	\
0		False	False	
1		False	False	
2		False	False	
3		False	False	
4		False	False	

	StreamingMovies_No	internet service	StreamingMovies_Yes
0		False	False
1		False	False
2		False	False
3		False	False
4		False	False

[5 rows x 31 columns]